Free Test

Completed on 17-December-2019

Attempt
01

Marks Obtained
9 / 15

Your score
60%

Time Taken
N/A

Result
Failed

Domains wise Quiz Performance Report

No	1
Domain	Processing
Total Question	5
Correct	5
Incorrect	0
Unattempted	0
Marked for review	0

No	2
Domain	Storage
Total Question	3
Correct	0
Incorrect	3
Unattempted	0
Marked for review	0

No	3
Domain	Collection
Total Question	1
Correct	1
Incorrect	0
Unattempted	0
Marked for review	0

No	4
Domain	Data Security
Total Question	4
Correct	2
Incorrect	2
Unattempted	0
Marked for review	0

No	5
Domain	Analysis
Total Question	1
Correct	1
Incorrect	0
Unattempted	0
Marked for review	0
No	6
Domain	Visualization
Total Question	1
Correct	0
Incorrect	1
Unattempted	0
Marked for review	0
Total	Total
All Domain	All Domain
Total Question	15
Correct	9
Incorrect	6
Unattempted	0
Marked for review	0

Review the Answers

Sorting by All

Question 1 **Correct**

Domain :Processing

Tick-Bank is a privately held Internet retailer of both physical and digital products founded in 2008. The company has more than six-million clients worldwide. Tick-Bank's technology aids in payments, tax calculations and a variety of customer service tasks and serve as a connection between digital content makers and affiliate dealers, who then promote them to clients thereby assist in building revenue making opportunities for companies.
Tick-Bank currently runs multiple java based web applications running on AWS and looking to enable web-site traffic analytics and also planning to extend the functionality for new web applications that are being launched. Tick-Bank uses KPL library to address event integration into the kinesis streams and thereby process the data to downstream applications for analytics. With growing applications and

customers, performance issues are hindering real time analytics and need an administrator to standardize performance, monitoring, manage and costs by kinesis streams.
Please select 3 options.

- ✓ A. **Use multiple shards to integrate data from different applications, reshard by splitting hot shards to increase capacity of the stream** ✓
- B. **Use multiple shards to integrate data from different applications, reshard by splitting cold shards to increase capacity of the stream**
- ✓ C. **Use CloudWatch metrics to monitor and determine the "hot" or "cold" shards and understand the usage capacity** ✓
- ✓ D. **Use multiple shards to integrate data from different applications, reshard by merging cold shards to reduce cost of the stream** ✓
- E. **Use multiple shards to integrate data from different applications, reshard by merging hot shards to reduce cost of the stream and improve performance**
- F. **Use CloudTrail metrics to monitor and determine the "hot" or "cold" shards and understand the usage capacity**

Explanation:

Answer: A, C, D

Option A is correct - Splitting hot shards improve the performance. Define a single shard for each web application of a kinesis stream. Based on the usage generated through CloudWatch Metrics for each shard, split the hot shards or merge the cold shards. This way we can improve the performance and reduce the costs for each stream. if the performance is still not addressed, adapt to different streams per application

https://docs.aws.amazon.com/streams/latest/dev/kinesis-using-sdk-java-resharding-strategies.html

Option B is incorrect - Splitting cold shards does not improve the performance. Define a single shard for each web application of a kinesis stream. Based on the usage generated through CloudWatch Metrics for each shard, split the hot shards or merge the cold shards. This way we can improve the performance and reduce the costs for each stream. if the performance is still not addressed, adapt to different streams per application

https://docs.aws.amazon.com/streams/latest/dev/kinesis-using-sdk-java-resharding-strategies.html

Option C is correct - CloudWatch Metrics determine which are your "hot" or "cold" shards, that is, shards that are receiving much more data, or much less data, than expected. You could then selectively split the hot shards to increase capacity for the hash keys that target those shards. Similarly, you could merge cold shards to make better use of their unused capacity.

https://docs.aws.amazon.com/streams/latest/dev/kinesis-using-sdk-java-resharding-strategies.html

Option D is correct - Merging cold shards improve performance as well as costs. Define a single shard for each web application of a kinesis stream. Based on the usage generated through CloudWatch Metrics for each shard, split the hot shards or merge the cold shards. This way we can improve the performance and reduce the costs for each stream. if the performance is still not addressed, adapt to different streams per application

https://docs.aws.amazon.com/streams/latest/dev/kinesis-using-sdk-java-resharding-strategies.html

Option E is incorrect - Merging hot shards does not improve performance. Define a single shard for each web application of a kinesis stream. Based on the usage generated through CloudWatch Metrics for each shard, split the hot shards or merge the cold shards. This way we can improve the performance and reduce the costs for each stream. if the performance is still not addressed, adapt to different streams per application

https://docs.aws.amazon.com/streams/latest/dev/kinesis-using-sdk-java-resharding-strategies.html

Option F is incorrect - CloudTrail does not provide this. CloudWatch Metrics determine which are your "hot" or "cold" shards, that is, shards that are receiving much more data, or much less data, than expected. You could then selectively split the hot shards to increase capacity for the hash keys that target those shards. Similarly, you could merge cold shards to make better use of their unused capacity.

https://docs.aws.amazon.com/streams/latest/dev/kinesis-using-sdk-java-resharding-strategies.html

Ask our Experts

Rate this Question? ☺ ☹

Question 2 Incorrect

Domain :Storage

Hymutabs Ltd (Hymutabs) is a global environmental solutions company running its operations in in Asia Pacific, the Middle East, Africa and the Americas. It maintains more than 10 exploration labs around the world, including a knowledge centre, an "innovative process development centre" in Singapore, a materials and membrane products development centre as well as advanced machining, prototyping and industrial design functions.
Hymutabs hosts their existing enterprise infrastructure on AWS and runs multiple applications to address the product life cycle management.

The datasets are available in Aurora, RDS and S3 in file format. Hymutabs Management team is interested in building analytics around product life cycle and advanced machining, prototyping and other functions.
The IT team proposed Redshift to fulfill the EDW and analytics requirements. They adapt modeling approaches laid by Bill Inmon and Kimball to efficiently design the solution. The team understands that the data loaded into Redshift would be in terabytes and identified multiple massive dimensions, facts, summaries of millions of records and are working on establishing the best practices to address the design concerns.
There are 6 tables that they are currently working on:

ORDER_FCT is a Fact Table with billions of rows related to orders

SALES_FCT is a Fact Table with billions of rows related to sales transactions. This table is specifically used to generate reports EOD (End of Day), EOW(End of Week), and EOM (End of Month) and also sales queries

?CUST_DIM is a Dimension table with billions of rows related to customers. It is a TYPE 2 Dimension table

PART_DIM is a part dimension table with billions of records that defines the materials that were ordered

DATE_DIM is a dimension table

SUPPLIER_DIM holds the information about suppliers the Hymutabs work with

One of the key requirements includes ORDER_FCT and PART_DIM are joined together in most of order related queries. ORDER_FCT has many other dimensions to support analysis.
How would you design the distribution? Select 1 option.

A. **Distribute the ORDER_FCT with KEY distribution on its primary KEY (any one of the columns) and PART_DIM with KEY distribution on its PRIMARY KEY**

✓ B. **Distribute the ORDER_FCT with ALL distribution on its primary KEY (any one of the columns) and PART_DIM with ALL distribution on its PRIMARY KEY** ✗

C. **Distribute the ORDER_FCT with EVEN distribution on its primary KEY (any one of the columns) and PART_DIM with EVEN distribution on its PRIMARY KEY**

D. **Distribute the ORDER_FCT and PART_DIM on same key with KEY distribution** ✓

E. **Distribute the ORDER_FCT and PART_DIM on same key with EVEN distribution**

Explanation:

Answer : D

Option A is incorrect - KEY DISTRIBUTION distributes the rows are according to the values in one column. Queries initiate lot of redistribution of data of both ORDER_FCT and PART_DIM are not built on same key.

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option B is incorrect - ALL distribution makes a copy of the entire table in every compute node. Being billion record tables, this is not a right approach to design.

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option C is incorrect - EVEN DISTRIBUTION evenly distributes the rows across the slices in a round-robin fashion, regardless of the values in any particular column. EVEN distribution is appropriate when a table does not participate in joins. Definitely not a right approach.

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option D is correct - KEY DISTRIBUTION distributes the rows are according to the values in one column. With distribution of data on same key in both the tables, there is no change of redistribution. This is the best approach to design.

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option E is incorrect - EVEN DISTRIBUTION evenly distributes the rows across the slices in a round-robin fashion, regardless of the values in any

particular column. EVEN distribution is appropriate when a table does not participate in joins. Definitely not a right approach
https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Ask our Experts

Rate this Question? ☺ ☹

Question 3 Correct

Domain :Processing

MSP Bank, Limited is a leading varied Japanese monetary institution that provides a full range of financial products and services to both institutional and individual customers. It is headquartered in Tokyo. MSP Bank is hosting their existing infrastructure on AWS. MSP bank has many segments internally and they are planning to launch a self-data discovery platform running out of AWS on QuickSight.
Using QuickSight, multiple datasets are created and multiple analyses are generated respectively. The Team is working on enabling auditing to track the records of actions taken by a user, role, or an AWS service in Amazon QuickSight. Also the team need to capture the logs and storage it for long term archival to address compliance. Please advice. Select 3 options.

- ✓ A. **Amazon QuickSight is integrated with AWS CloudTrail which provides a record of actions taken by a user, role, or an AWS service in Amazon QuickSight** ✓
- B. **Amazon QuickSight is integrated with AWS CloudWatch which provides a record of actions taken by a user, role, or an AWS service in Amazon QuickSight**
- ✓ C. **when CloudTrail is enabled, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Amazon QuickSight** ✓
- D. **when CloudWatch is enabled, you can enable continuous delivery of CloudWatch events to an Amazon S3 bucket, including events for Amazon QuickSight**
- ✓ E. **If you don't configure a trail, you can still view the most recent events in the CloudTrail console in Event history** ✓
- F. **If you don't configure a log, you can still view the most recent events in the CloudWatch console in Event history**

Explanation:

Answer: A,C,E
Amazon QuickSight is integrated with AWS CloudTrail. This service provides a record of actions taken by a user, role, or an AWS service in Amazon QuickSight. The calls captured include calls from the Amazon QuickSight console. If you create a trail, you can enable continuous delivery of CloudTrail events to an Amazon S3 bucket, including events for Amazon QuickSight. If you don't configure a trail, you can still view the most recent events in the CloudTrail console in Event history. Using the information collected by CloudTrail, you can determine the request that was made to Amazon QuickSight, the IP address from which the request was made, who made the request, when it was made, and additional details
https://docs.aws.amazon.com/quicksight/latest/user/logging-using-cloudtrail.html

Ask our Experts

Rate this Question?

Question 4 Incorrect

Domain :Storage

ConsumersHalt (CH) is an Indian department collection chain. There are 63 branches across 32 towns in India, with clothing, accessories, bags, shoes, jewelry, scents, faces, health and exquisiteness products, home furnishing and decor products.
CH runs their existing operations and analytics infrastructure out of AWS which includes S3, EC2, Auto Scaling, CDN and also Redshift. The Redshift platform is being used for advanced analytics, real time analytics and being actively used for past 2 years. Suddenly performance issues are occurring in the application and administrator being a superuser needs to provide a list of reports in terms of current and historical performance of the cluster. What types of tables/views can help access the performance related info for diagnosis.
Select 3 options.

- ✓ A. **STL system tables are generated from Amazon Redshift log files to provide a history of the system. They serve logging.** ✓
- B. **STL tables are actually virtual system tables that contain snapshots of the current system data. They serve snapshots.**
- ✓ C. **STV system tables are generated from Amazon Redshift log files to provide a history of the system. They serve logging.** ✗
- D. **STV tables are actually virtual system tables that contain snapshots of the current system data. They serve snapshots.** ✓
- ✓ E. **System views contain full data found in several of the STL and STV system tables.** ✗
- F. **The system catalogs store schema metadata, such as information about tables and columns.** ✓

Explanation:

Answer : A, D, F

Option A is correct - STL system tables are generated from Amazon Redshift log files to provide a history of the system.

https://docs.aws.amazon.com/redshift/latest/dg/c_intro_STL_tables.html

Option B is incorrect - STL system tables are generated from Amazon Redshift log files to provide a history of the system.

https://docs.aws.amazon.com/redshift/latest/dg/c_intro_STL_tables.html

Option C is incorrect - STV tables are actually virtual system tables that contain snapshots of the current system data.

https://docs.aws.amazon.com/redshift/latest/dg/c_intro_STV_tables.html

Option D is correct - STV tables are actually virtual system tables that contain snapshots of the current system data.

https://docs.aws.amazon.com/redshift/latest/dg/c_intro_STV_tables.html

Option E is incorrect - System tables contain only subset of data

https://docs.aws.amazon.com/redshift/latest/dg/c_intro_system_views.html

Option F is correct - The system catalogs store schema metadata, such as information about tables and columns. System catalog tables have a PG prefix.

https://docs.aws.amazon.com/redshift/latest/dg/c_intro_catalog_views.html

Ask our Experts

Rate this Question? ☺ ☹

Question 5 Correct

Domain :Processing

HikeHills.com (HH) is an online specialty retailer that sells clothing and outdoor refreshment gear for trekking, go camping, boulevard biking, mountain biking, rock hiking, ice mountaineering, skiing, avalanche protection, snowboarding, fly fishing, kayaking, rafting, road and trace running, and many more.
HH runs their entire online infrastructure on multiple java based web applications and other web framework applications running on AWS. The HH is capturing clickstream data and use custom-build recommendation engine to recommend products which eventually improve sales, understand customer preferences and already using AWS Kinesis Streams (KDS) to collect events and transaction logs and process the stream. Multiple departments from HH use different streams to address real-time

integration and induce analytics into their applications and uses Kinesis as the backbone of real-time data integration across the enterprise.
HH uses a VPC to host all their applications and is looking at integration of kinesis into their web application. To understand the network flow behavior based on every 15 minutes, HH is looking at aggregating data based on the VPC logs for analytics. VPC Flow Logs have a capture window of approximately 10 minutes. What kind of queries can be used to capture aggregates based on each client for every 15 mins using Amazon Kinesis Data Analytics. Select 1 option.

- ✓ A. **Stagger Windows queries** ✅
- B. **Tumbling Windows queries**
- C. **Sliding windows queries**
- D. **Continuous queries**

Explanation:

Answer: A

Option A is correct - Stagger windows query, A query that aggregates data using keyed time-based windows that open as data arrives. The keys allow for multiple overlapping windows. This is the recommended way to aggregate data using time-based windows.

VPC Flow Logs have a capture window of approximately 10 minutes. But they can have a capture window of up to 15 minutes if you're aggregating data on the client. Stagger windows are ideal for aggregating these logs for analysis.
https://docs.aws.amazon.com/kinesisanalytics/latest/dev/stagger-window-concepts.html

Option B is incorrect - Tumbling Windows query, A query that aggregates data using distinct time-based windows that open and close at regular intervals.

https://docs.aws.amazon.com/kinesisanalytics/latest/dev/tumbling-window-concepts.html

Option C is incorrect - Sliding windows query, A query that aggregates data continuously, using a fixed time or rowcount interval.

https://docs.aws.amazon.com/kinesisanalytics/latest/dev/sliding-window-concepts.html

Option D is incorrect - Continuous Query is a query over a stream executes continuously over streaming data. This continuous execution enables scenarios, such as the ability for applications to continuously query a stream and generate alerts.

https://docs.aws.amazon.com/kinesisanalytics/latest/dev/continuous-queries-concepts.html

Ask our Experts

Rate this Question?

Question 6 Correct

Domain :Processing

HikeHills.com (HH) is an online specialty retailer that sells clothing and outdoor refreshment gear for trekking, go camping, boulevard biking, mountain biking, rock hiking, ice mountaineering, skiing, avalanche protection, snowboarding, fly fishing, kayaking, rafting, road and trace running, and many more.
HH runs their entire online infrastructure on java based web applications running on AWS. The HH is capturing click stream data and use custom-build recommendation engine to recommend products which eventually improve sales, understand customer preferences and already using AWS kinesis KPL to collect events and transaction logs and process the stream. The event/log size is around 12 bytes. The log stream generated into the stream is used for multiple purposes. HH proposes Kinesis Firehose to process the stream and capture information. What purposes can be fulfilled OOTB without writing applications or consumer code? Select 4 options.

- ✓ A. **Deliver real-time streaming data to Amazon Simple Storage Service (Amazon S3)** ✓
- B. **Deliver real-time streaming data to DynamoDB to support processing of digital documents**
- ✓ C. **Deliver real-time streaming data to Redshift to support data warehousing and real-time analytics** ✓
- ✓ D. **Ingest data into ES domains to support Enterprise search built on Elasticsearch** ✓
- ✓ E. **Allow Splunk to read and process data stream directly from Kinesis Firehose** ✓
- F. **Ingest data into Amazon EMR to support big data analytics**

Explanation:

Answer: A, C, D, E

Amazon Kinesis Data Firehose is a fully managed service for delivering real-time streaming data to destinations such as Amazon Simple Storage Service (Amazon S3), Amazon Redshift, Amazon Elasticsearch Service (Amazon ES), and Splunk. With Kinesis Data Firehose, you don't need to write applications or manage resources. Configure data producers to send data to Kinesis Data Firehose,

and it automatically delivers the data to the destination that you specified. You can also configure Kinesis Data Firehose to transform your data before delivering it.

https://docs.aws.amazon.com/firehose/latest/dev/what-is-this-service.html

Ask our Experts

Rate this Question? ☺ ☹

Question 7 Incorrect

Domain :Storage

Hymutabs Ltd (Hymutabs) is a global environmental solutions company running its operations in in Asia Pacific, the Middle East, Africa and the Americas. It maintains more than 10 exploration labs around the world, including a knowledge centre, an "innovative process development centre" in Singapore, a materials and membrane products development centre as well as advanced machining, prototyping and industrial design functions.
Hymutabs hosts their existing enterprise infrastructure on AWS and runs multiple applications to address the product life cycle management. The datasets are available in Aurora, RDS and S3 in file format. Hymutabs Management team is interested in building analytics around product life cycle and advanced machining, prototyping and other functions.
The IT team proposed Redshift to fulfill the EDW and analytics requirements. They adapt modeling approaches laid by Bill Inmon and Kimball to efficiently design the solution. The team understands that the data loaded into Redshift would be in terabytes and identified multiple massive dimensions, facts, summaries of millions of records and are working on establishing the best practices to address the design concerns.
There are 6 tables that they are currently working on:

ORDER_FCT is a Fact Table with billions of rows related to orders

SALES_FCT is a Fact Table with billions of rows related to sales transactions. This table is specifically used to generate reports EOD (End of Day), EOW(End of Week), and EOM (End of Month) and also sales queries

CUST_DIM is a Dimension table with billions of rows related to customers. It is a TYPE 2 Dimension table

PART_DIM is a part dimension table with billions of records that defines the materials that were ordered

DATE_DIM is a dimension table

SUPPLIER_DIM holds the information about suppliers the Hymutabs work with

SALES_FCT and DATE_DIM are joined together frequently since EOD sales reports are generated every day. please suggest your distribution style for both tables. Select 1 option.

A. Distribute the SALES_FCT with KEY DISTRIBUTION on its own Primary KEY (one of the columns) while DATE_DIM is distributed with KEY DISTRIBUTION on Its PRIMARY KEY

B. Distribute the SALES_FCT with EVEN DISTRIBUTION on its own Primary KEY (one of the columns) while DATE_DIM is distributed with EVEN distribution on Its PRIMARY KEY

C. Distribute the SALES_FCT with KEY DISTRIBUTION on its own Primary KEY (one of the columns) while DATE_DIM is distributed with ALL DISTRIBUTION on Its PRIMARY KEY

D. Distribute the SALES_FCT with ALL DISTRIBUTION on its own Primary KEY (one of the columns) while DATE_DIM is distributed with EVEN distribution on Its PRIMARY KEY

E. Distribute the SALES_FCT with EVEN DISTRIBUTION on its own Primary KEY (one of the columns) while DATE_DIM is distributed with ALL distribution on Its PRIMARY KEY

Explanation:

Answer : C

Option A is incorrect - KEY DISTRIBUTION distributes the rows are according to the values in one column. This is a right approach to design the table, but DATE_DIM with KEY DISTRIBUTION with number of records being very low, lot of data is copied between nodes. This approach is ok but not a perfect design to build the solution

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option B is incorrect - EVEN DISTRIBUTION evenly distributes the rows across the slices in a round-robin fashion, regardless of the values in any particular column. EVEN distribution is appropriate when a table does not participate in joins. For a fact table like SALES_FCT, all the nodes participate in all queries even though the EOD reports is only for that particular day

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option C is correct - ALL distribution makes a copy of the entire table in every compute node. Being billion record tables, this is not a right approach to design. This is the perfect design for DATE_DIM table which has very low number and can be distributed to all tables

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option D is incorrect - ALL distribution makes a copy of the entire table in every compute node. Being billion record tables, this is not a right approach to design. Cannot be used for massive table

like SALES_FCT.

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Option E is incorrect - EVEN DISTRIBUTION evenly distributes the rows across the slices in a round-robin fashion, regardless of the values in any particular column. EVEN distribution is appropriate when a table does not participate in joins. For a fact table like SALES_FCT, all the nodes participate in all queries even though the EOD reports is only for that particular day. SALES_FCT TABLE need to be designed on a table with a perfect distribution key in mind

https://docs.aws.amazon.com/redshift/latest/dg/tutorial-tuning-tables-distribution.html

Ask our Experts

Rate this Question?

Question 8 Correct

Domain :Collection

Tick-Bank is a privately held Internet retailer of both physical and digital products founded in 2008. The company has more than six-million clients worldwide. Tick-Bank aims to serve as a connection between digital content makers and affiliate dealers, who then promote them to clients. Tick-Bank's technology aids in payments, tax calculations and a variety of customer service tasks. Tick-Bank assists in building perceptibility and revenue making opportunities for entrepreneurs.
Tick-Bank runs multiple java based web applications running on windows based EC2 machines in AWS managed by internal IT Java team, to serve various business functions. Tick-Bank is looking to enable web-site traffic analytics there by understanding user navigational behavior, preferences and other click related info. The amount of data captured per click is in tens of bytes. Tick-Bank has the following objectives in mind for the solution.
Tick-Bank uses KPL to process the data and KCL library to consume the records. Thousands of events are being generated every second and every event is sensitive and equally important and Gluebush.com wants to treat every record as a separate stream. please detail the implementation guidelines. select 2 options.

- ✓ A. **each record in a separate Kinesis Data Streams record and make one HTTP request to send it to Kinesis Data Streams** ✓
- B. **each HTTP request carries multiple Kinesis Stream records which is sent to kinesis Data streams**
- C. **Batching is implemented as the target implementation**
- ✓ D. **Batching is not implemented as the target implementation** ✓

Explanation:

Answer: A,D

Option A is correct - When batching is not implemented, each record is treated as a separate data stream record. In this context batching is not implemented.

https://docs.aws.amazon.com/streams/latest/dev/kinesis-kpl-concepts.html

Option B is incorrect - When batching is not implemented, each record is treated as a separate data stream record. In this context batching is not implemented.

https://docs.aws.amazon.com/streams/latest/dev/kinesis-kpl-concepts.html

Option C is incorrect - Batching refers to performing a single action on multiple items instead of repeatedly performing the action on each individual item. In this context, the "item" is a record, and the action is sending it to Kinesis Data Streams. In a non-batchingsituation, you would place each record in a separate Kinesis Data Streams record and make one HTTP request to send it to Kinesis Data Streams. With batching, each HTTP request can carry multiple records instead of just one.

https://docs.aws.amazon.com/streams/latest/dev/kinesis-kpl-concepts.html

Option D is correct - Batching refers to performing a single action on multiple items instead of repeatedly performing the action on each individual item. In this context, the "item" is a record, and the action is sending it to Kinesis Data Streams. In a non-batchingsituation, you would place each record in a separate Kinesis Data Streams record and make one HTTP request to send it to Kinesis Data Streams. With batching, each HTTP request can carry multiple records instead of just one.

https://docs.aws.amazon.com/streams/latest/dev/kinesis-kpl-concepts.html

Ask our Experts

Rate this Question? ☺ ☹

Question 9 Incorrect

Domain :Data Security

Allianz Financial Services (AFS) is a banking group offering end-to-end banking and financial solutions in South East Asia through its consumer banking, business banking, Islamic banking, investment finance and stock broking businesses as well as unit trust and asset administration, having served the financial community over the past five decades.

AFS uses Redshift on AWS to fulfill the data warehousing needs and uses S3 as the staging area to host files. AFS uses other services like DynamoDB, Aurora, and Amazon RDS on remote hosts to fulfill other needs. AFS want to implement Redshift security end to end. How can this be achieved? select 2 options.

A. **Access to your Amazon Redshift Management Console is controlled by your AWS account privileges** ✓

B. **Define a cluster security group and associate it with a cluster to control access to specific Amazon Redshift resources** ✓

✓ C. **To encrypt the connection between your SQL client and your cluster, enable cluster encryption when you launch the cluster** ✗

✓ D. **To encrypt the data in all your user-created tables, you can use secure sockets layer (SSL) encryption** ✗

Explanation:

Answer : A,B

Amazon Redshift database security is distinct from other types of Amazon Redshift security. In addition to database security, which is described in this section, Amazon Redshift provides these features to manage security:

Sign-in credentials — Access to your Amazon Redshift Management Console is controlled by your AWS account privileges. For more information, see Sign-In Credentials.

Access management — To control access to specific Amazon Redshift resources, you define AWS Identity and Access Management (IAM) accounts. For more information, see Controlling Access to Amazon Redshift Resources.

Cluster security groups — To grant other users inbound access to an Amazon Redshift cluster, you define a cluster security group and associate it with a cluster. For more information, see Amazon Redshift Cluster Security Groups.

VPC — To protect access to your cluster by using a virtual networking environment, you can launch your cluster in an Amazon Virtual Private Cloud (VPC). For more information, see Managing Clusters in Virtual Private Cloud(VPC).

Cluster encryption — To encrypt the data in all your user-created tables, you can enable cluster encryption when you launch the cluster. For more information, see Amazon Redshift Clusters.

SSL connections — To encrypt the connection between your SQL client and your cluster, you can use secure sockets layer (SSL) encryption. For more information, see Connect to Your Cluster Using SSL.

Load data encryption — To encrypt your table load data files when you upload them to Amazon S3, you can use either server-side encryption or client-sideencryption. When you load from server-side encrypted data, Amazon S3 handles decryption transparently. When you load from client-side encrypted data, the Amazon Redshift COPY command decrypts the data as it loads the table. For more information, see Uploading Encrypted Data to Amazon S3.

Data in transit — To protect your data in transit within the AWS cloud, Amazon Redshift uses hardware accelerated SSL to communicate with Amazon S3 or Amazon DynamoDB for COPY, UNLOAD, backup, and restore operations.

https://docs.aws.amazon.com/redshift/latest/dg/c_security-overview.html

Ask our Experts

Rate this Question?

Question 10 Correct

Domain :Processing

Parson Fortunes Ltd is an Asian-based department store operator with an extensive network of 131 stores, spanning approximately 4.1 million m2 of retail space across cities in India, China, Vietnam, Indonesia, and Myanmar.
Parson has large assets of data around 10 TB's of structured data and 5 TB of unstructured data and is planning to host their data warehouse on AWS and unstructured data storage on S3. Parson IT team is well aware of the scalability, performance of AWS services capabilities. Parson is currently using running their DWH, on-premises on Teradata and is concerned on the overall costs of the DWH on AWS. They want to initially migrate the platform onto AWS use it for basic analytics, and don't have any performance intensive workloads in place for time being. They have business needs around real-time data integration, data driven analytics as a roadmap of 5 years. Currently the number of users accessing the application would be around 100. What is your suggestion? select 1 option.

✓ A. **Launch Redshift cluster with node types DS2.xlarge to fulfill the requirements**

B. **Launch Redshift cluster with node types DS2.8xlarge to fulfill the requirements**

C. **Launch Redshift cluster with node types DC2.xlarge to fulfill the requirements**

D. **Launch Redshift cluster with node types DC2.8xlarge to fulfill the requirements**

Explanation:

Answer : A

Option A is correct - DS2 node types are optimized for large data workloads and use hard disk drive (HDD) storage.DS2.xlarge fulfills the requirements since it provides massive parallel processing using multiple nodes. Based on the amount of data loaded, this is the right option

https://docs.aws.amazon.com/redshift/latest/mgmt/working-with-clusters.html#rs-about-clusters-and- nodes

Option B is INCORRECT ,since , although cost-wise both the ds2.xlarge and ds2.8xlarge are the same, deploying a ds2.xlarge instance would be just sufficient as per the requirements of the question

Please refer the following link
https://aws.amazon.com/blogs/aws/amazon-redshift-now-faster-and-more-cost-effective-than-ever/

Option C is incorrect - DC2 node types are optimized for performance-intensive workloads. This is not the requirement

https://docs.aws.amazon.com/redshift/latest/mgmt/working-with-clusters.html#rs-about-clusters-and- nodes

Option D is incorrect - DC2 node types are optimized for performance-intensive workloads. This is not the requirement

https://docs.aws.amazon.com/redshift/latest/mgmt/working-with-clusters.html#rs-about-clusters-and- nodes

Ask our Experts

Rate this Question? ☺ ☹

Question 11 Correct

Domain :Data Security

QuickDialog is a multimedia company running a messaging app. One of the principal features of QuickDialog is that pictures and messages are usually only available for a short time before they become inaccessible to users. The app has evolved from originally centering on person-to-person photo sharing to present users' "Stories" of 24 hours of sequential content, along with "Discover", allowing brands show ad-supported short-form media.

They use DynamoDB to support the mobile application and S3 to host the images and other documents shared between users. KindleYou has a large customer base spread across multiple geographic areas. Customers need to update their profile information while using the application. Propose a solution that can be easily implemented and provides full consistency.
Select 1 Option

- A. **Use global tables, a fully managed solution across multiple regions, multi-master databases** ✓
- B. **Create CustomerProfile table in a region, create replication copies in different AWS regions and enable replication through AWS Kinesis Data Streams**
- C. **Create CustomerProfile table in a region, create replication copies in different AWS regions and enable replication through AWS Data Pipeline**
- D. **Create CustomerProfile table in a region, create replication copies in different AWS regions and enable replication through AWS Kinesis Data Firehose**

Explanation:

Answer: A

Option A is correct - Amazon DynamoDB global tables provide a fully managed solution for deploying a multi- region, multi-master database, without having to build and maintain your own replication solution. Replication is performed by DynamoDB streams

https://docs.aws.amazon.com/amazondynamodb/latest/developerguide/GlobalTables.html

Option B is incorrect - Replication with Kinesis Streams is possible, but does not provide full consistency. But again producers using KPL and Consumers need to be configured

https://docs.aws.amazon.com/streams/latest/dev/key-concepts.html

Option C is incorrect - Replication with Data Pipeline is based on pipelines and schedule. It cannot be real-time and does not provide full consistency

https://docs.aws.amazon.com/datapipeline/latest/DeveloperGuide/what-is-datapipeline.html

Option D is incorrect - Replication with Kinesis Firehose runs with a lag. It cannot be real-time and does not provide full consistency

https://docs.aws.amazon.com/firehose/latest/dev/what-is-this-service.html

Ask our Experts

Rate this Question?

Question 12 Incorrect

Domain :Data Security

HikeHills.com (HH) is an online specialty retailer that sells clothing and outdoor refreshment gear for trekking, go camping, boulevard biking, mountain biking, rock hiking, ice mountaineering, skiing, avalanche protection, snowboarding, fly fishing, kayaking, rafting, road and trace running, and many more.
HHruns their entire online infrastructure on java based web applications running on AWS. The HH is capturing clickstream data and use custom-build recommendation engine to recommend products which eventually improve sales, understand customer preferences and already using AWS Streaming capabilities to collect events and transaction logs and process the stream.
HHis using kinesis analytics to build SQL querying capability on streaming and planning to use different types of queries to process the data. HH need to ensure proper authentication and authorization control for kinesis analytics application needs to be enabled. How can this be achieved? select 2 options.

- ✓ A. **Authentication and Access to AWS resources using following identities like root user, IAM User, and IAM role thereby managing federated user access, AWS service access and Applications running on Amazon EC2** ✓
- B. **Access Control using following identities like root user, IAM User, and IAM role thereby managing federated user access, AWS service access and Applications running on Amazon EC2**
- ✓ C. **Authentication and Access to AWS resources through Permissions, policies, Actions and Resources** ✗
- D. **Access Control through Permissions, policies, Actions and Resources** ✓

Explanation:

Answer: A, D

Access to Amazon Kinesis Data Analytics requires credentials. Those credentials must have permissions to access AWS resources, such as an Amazon Kinesis Data Analytics application or an Amazon Elastic Compute Cloud (Amazon EC2) instance.

Authentication - root user, IAM User, and IAM role thereby managing federated user access, AWS service access and Applications running on Amazon EC2

Access Control - through Permissions, policies, Actions and Resources

https://docs.aws.amazon.com/kinesisanalytics/latest/dev/authentication-and-access-control.html

Ask our Experts

Rate this Question?

Question 13 Correct

Domain :Data Security

As a part of the smart city initiatives, Hyderabad (GHMC), one of the largest cities in southern India is working on capturing massive volumes of video streams 24/7 captured from the large numbers of "Vivotek IB9371 – HT" cameras installed at traffic lights, parking lots, shopping malls, and just about every public venue to help solve traffic problems, help prevent crime, dispatch emergency responders, and much more. GHMC uses AWS to host their entire infrastructure.
The camera's write stream into Kinesis Video Stream securely and eventually consumed by applications for custom video processing, on-demand video playback and also consumed by AWS Rekognition for video analytics. Along with the stream, different modes of streaming metadata are sent along with the stream. There are 2 scenarios that need to be fulfilled.
·Requirement 1 - Affix metadata on a specific Adhoc basis to fragments in a stream, aka when smart camera detects motion in restricted areas, adds metadata [Motion = true] to the corresponding fragments that contain the motion before sending the fragments to its Kinesis Video Stream
·Requirement 2 - affix metadata to successive, consecutive fragments in a stream based on a continuing need, aka all smart cameras in the city sends the current latitude and longitude coordinates associated with all fragments it sends to its Kinesis Video Stream
How can this be achieved? Select 2 options.

- ✓ A. **Requirement 1 can be fulfilled by sending Nonpersistent data** ✓
- B. **Requirement 2 can be fulfilled by sending Nonpersistent dat**
- C. **Requirement 1 can be fulfilled by sending Persistent data**
- ✓ D. **Requirement 2 can be fulfilled by sending Persistent data** ✓
- E. **Both Requirement 1 and Requirement 2 can be fulfilled by sending Nonpersistent data**
- F. **Both Requirement 1 and Requirement 2 can be fulfilled by sending Persistent data**

Explanation:

Answer: A, D

Option A is correct. Requirement 1 can be fulfilled by sending Nonpersistent data.

https://docs.aws.amazon.com/kinesisvideostreams/latest/dg/how-meta.html

Option B is incorrect. Requirement 2 cannot be fulfilled by sending Nonpersistent data.

https://docs.aws.amazon.com/kinesisvideostreams/latest/dg/how-meta.html

Option C is incorrect. Requirement 1 cannot be fulfilled by sending Persistent data.

https://docs.aws.amazon.com/kinesisvideostreams/latest/dg/how-meta.html

Option D is correct. Requirement 2 can be fulfilled by sending Persistent data.

https://docs.aws.amazon.com/kinesisvideostreams/latest/dg/how-meta.html

Option E is incorrect. Requirement 1 can be fulfilled by sending Nonpersistent data while Requirement 2 can be fulfilled by sending Persistent data.

https://docs.aws.amazon.com/kinesisvideostreams/latest/dg/how-meta.html

O**ption F is incorrect.** Requirement 1 cannot be fulfilled by sending Persistent data while Requirement 2 can be fulfilled by sending Persistent data.

https://docs.aws.amazon.com/kinesisvideostreams/latest/dg/how-meta.html

Ask our Experts

Rate this Question? ☺ ☹

Question 14 Correct

Domain :Analysis

Marqueguard is a social media monitoring company headquartered in Brighton, England. Marqueguard sells three different products: Analytics, Audiences, and Insights. Marqueguard Analytics is a "self-serveapplication" or software as a service, which archives social media data in order to provide companies with information and the means to track specific segments to analyze their brands' online presence.
The tool's coverage includes blogs, news sites, forums, videos, reviews, images and social networks such as Twitter and Facebook. Users can search data by using Text and Image Search, and use

charting, categorization, sentiment analysis and other features to provide further information and analysis. Marqueguard has access to over 80 million sources.
Marqueguard wants provide Image and text analysis capabilities to the applications which includes identify objects, people, text, scenes, and activities and also also provides highly accurate facial analysis and facial recognition.
What service can provide this capability? select 1 option.

A. **Amazon Comprehend**

✓ B. **Amazon Rekognition**

C. **Amazon Polly**

D. **Amazon SageMaker**

Explanation:

Answer : B

Option A is incorrect - Amazon Comprehend uses natural language processing (NLP) to extract insights about the content of documents. Amazon Comprehend processes any text file in UTF-8 format. It develops insights by recognizing the entities, key phrases, language, sentiments, and other common elements in a document. Use Amazon Comprehend to create new products based on understanding the structure of documents. For example, using Amazon Comprehend you can search social networking feeds for mentions of products or scan an entire document repository for key phrases.

https://docs.aws.amazon.com/comprehend/latest/dg/what-is.html

Option B is correct - Amazon Rekognition makes it easy to add image and video analysis to your applications. You just provide an image or video to the Rekognition API, and the service can identify objects, people, text, scenes, and activities. It can detect any inappropriate content as well. Amazon Rekognition also provides highly accurate facial analysis and facial recognition. You can detect, analyze, and compare faces for a wide variety of use cases, including user verification, cataloging, people counting, and public safety.

https://docs.aws.amazon.com/rekognition/latest/dg/what-is.html

Option C is incorrect - Amazon Polly is a cloud service that converts text into lifelike speech. You can use Amazon Polly to develop applications that increase engagement and accessibility. Amazon Polly supports multiple languages and includes a variety of lifelike voices, so you can build speech-enabled applications that work in multiple locations and use the ideal voice for your customers. With Amazon Polly, you only pay for the text you synthesize. You can also cache and replay Amazon Polly's generated speech at no additional cost.

https://docs.aws.amazon.com/polly/latest/dg/what-is.html

Option D is incorrect - Amazon SageMaker is a fully managed machine learning service. With Amazon SageMaker, data scientists and developers can quickly and easily build and train machine learning models, and then directly deploy them into a production-readyhosted environment. It provides an integrated Jupyter authoring notebook instance for easy access to your data sources for exploration and analysis, so you don't have to manage servers. It also provides common machine learning algorithms that are optimized to run efficiently against extremely large data in a distributed environment

https://docs.aws.amazon.com/sagemaker/latest/dg/whatis.html

Ask our Experts

Rate this Question?

Question 15 Incorrect

Domain :Visualization

MSP Bank, Limited is a leading Japanese monetary institution that provides a full range of financial products and services to both institutional and individual customers. It is headquartered in Tokyo. MSP Bank is hosting their existing infrastructure on AWS. MSP bank has many segments internally and they are planning to launch a self-data discovery platform running out of AWS on QuickSight.
Using QuickSight, multiple datasets are created and multiple analyses are generated respectively. The Team is working on visuals. They are looking for a customized table view of data. Please help. select 1 option.

A. **Tabular Reports**

B. **Heat Maps**

C. **Pie Chart**

D. **Tree Map**

Explanation:

Answer: A

Option A is correct - Use tabular reports to see a customized table view of your data.

To create a table visual, choose at least one field of any data type. You can add as many columns as you need. Plus, you can add calculated columns.

https://docs.aws.amazon.com/quicksight/latest/user/tabular.html

Option B is incorrect - Use heat maps to show a measure for the intersection of two dimensions, with color- coding to easily differentiate where values fall in the range. Heat maps can also be used to show the count of values for the intersection of the two dimensions.

https://docs.aws.amazon.com/quicksight/latest/user/heat-map.html

Option C is incorrect - Use pie charts to compare values for items in a dimension.

Each wedge in a pie chart represents one item in the dimension. Wedge size represents the proportion of the value for the selected measure that the item represents compared to the whole for the dimension. Pie charts are best when precision isn't important and there are few items in the dimension.

https://docs.aws.amazon.com/quicksight/latest/user/pie-chart.html

Option D is incorrect - Use tree maps to visualize one or two measures for a dimension.

Each rectangle on the tree map represents one item in the dimension. Rectangle size represents the proportion of the value for the selected measure that the item represents compared to the whole for the dimension. You can optionally use rectangle color to represent another measure for the item. Rectangle color represents where the value for the item falls in the range for the measure, with darker colors indicating higher values and lighter colors indicating lower ones.

https://docs.aws.amazon.com/quicksight/latest/user/tree-map.html

Ask our Experts

Rate this Question?

Finish Review

Certification

Cloud Certification

Java Certification

PM Certification

Big Data Certification

Company

Support

Discussions

Blog

Business

Follow us