



Sparse arrays





About sparse arrays

- A **sparse array** is simply an array most of whose entries are zero (or **null**, or some other default value)
- For example: Suppose you wanted a 2-dimensional array of course grades, whose rows are Penn students and whose columns are courses
 - There are about 22,000 students
 - There are about 5000 courses
 - This array would have about 110,000,000 entries
 - Since most students take fewer than 5000 courses, there will be a lot of empty spaces in this array
 - This is a big array, even by modern standards
- There *are* ways to represent sparse arrays efficiently


Sparse arrays as linked lists

- We will start with sparse one-dimensional arrays, which are simpler
 - We'll do sparse two-dimensional arrays later
- Here is an example of a sparse one-dimensional array:

ary

0	1	2	3	4	5	6	7	8	9	10	11
0	0	0	0	17	0	0	23	14	0	0	0

- Here is how it could be represented as a linked list:





A sparse array ADT

- For a one-dimensional array of Objects, you would need:
 - A constructor: `SparseArray(int length)`
 - A way to get values from the array:
`Object fetch(int index)`
 - A way to store values in the array:
`void store(int index, Object value)`
- Note that it is *OK* to ask for a value from an “empty” array position
 - For an array of numbers, this should return zero
 - For an array of Objects, this should return `null`
- Additional useful operations:
 - `int length()` : return the size of the array
 - `int elementCount()` : how many non-null values are in the array
- Are there any important operations we have forgotten?



Implementing the operations

- ```
class List { int index; // the row number
 Object value; // the actual data
 List next; // the "pointer"

 public Object fetch(int index) {
 List current = this; // first "List" (node) in the list
 do {
 if (index == current.index) {
 return current.value; // found correct location
 }
 current = current.next;
 } while (index < current.index && next != null);
 return null; // if we get here, it wasn't in the list
 }
}
```
- The **store** operation is basically the same, with the extra complication that we may need to insert a node into the list



# Time analysis

---

- We must search a linked list for a given index
- We can keep the elements in order by index
- Expected time for both **fetch** and **store** is  $1/2$  the number of (nonzero/nonnull) elements in the list
- That is,  **$O(n)$** , where  $n$  is the number of actual (non-default) elements
  - For a “normal” array, indexing takes constant time
  - But for a sparse array, this isn’t bad
  - This is a typical example of a **time-space tradeoff**--in order to use less of one (space), we need to use more of the other (time)
- Expected time for the secondary methods, **length** and **elementCount**, is just  **$O(1)$** , that is, constant time
- We’re done, right?
- Unfortunately, this analysis is *correct but misleading*



# What is the problem?

---

- **True fact:** In an ordinary array, indexing to find an element is the only operation we really need
- **True fact:** In a sparse array, we can do indexing reasonably quickly
- **False conclusion:** In a sparse array, indexing to find an element is the only operation we really need
- The problem is that in designing the ADT, we didn't think enough about how it would be *used*



# Example: Finding the maximum

- To find the maximum element in a normal array:
  - ```
double max = array[0];  
for (int i = 0; i < array.length; i++) {  
    if (array[i] > max) max = array[i];  
}
```
- To find the maximum element in a sparse array:
 - ```
Double max = (Double) array.fetch(0);
for (int i = 0; i < array.length(); i++) {
 Double temp = (Double) array.fetch(i);
 if (temp.compareTo(max) > 0) {
 max = temp;
 }
}
```
- Do you see any problems with this?





# Problems with this approach

---

- Since we tried to be general, we defined our sparse array to hold Objects
  - This means a lot of wrapping and casting, which is awkward
  - We can deal with this (especially if we use Java 1.5)
- More importantly, in a normal array, every element is relevant
- If a sparse array is 1% full, 99% of its elements will be zero
  - This is 100 times as many elements as we should need to examine
- Our search time is based on the *size* of the sparse array, not on the number of elements that are actually in it
  - And it's a *big* array (else we wouldn't bother using a sparse array)



# Fixing the ADT

---

- Although “stepping through an array” is not a fundamental operation on an array, it is one we do frequently
  - Idiom: `for (int i = 0; i < array.length; i++) {...}`
- This is a very expensive thing to do with a sparse array
- This *shouldn't* be so expensive: We have a list, and all we need to do is step through it
  - *Poor* solution: Let the user step through the list
    - The user should not need to know anything about implementation
    - We cannot trust the user not to screw up the sparse array
    - These arguments are valid *even if the user is also the implementer!*
  - Correct solution: Expand the ADT by adding operations
    - But what, exactly, should these operations be?
    - Java *has* an answer, and it is the answer we should use



# Interfaces

---

- An **interface**, in Java, is like a class, but
  - It contains only **public** methods (and maybe some **final** values)
  - It only *declares* methods; it doesn't *define* them

- Example:

```
public interface Iterator { // Notice: no method bodies
 public boolean hasNext();
 public Object next();
 public void remove();
}
```

- This is an interface that is defined for you, in **java.util**
  - “Stepping through all the values” is something that you want to do for many data structures, so Java defines a standard way to do it
- You can write your own interfaces, using the above syntax
- So, how do you use this interface?



# Implementing an interface

---

- To use an interface, you say you are going to implement it, then you define every method in it
- Example:

```
public class SparseArrayIterator implements Iterator {
 // any data you need to keep track of goes here
 SparseArrayIterator() { ...an interface can't tell you what
 constructors to have, but you do need one... }
 public boolean hasNext () { ...you write this code... }
 public Object next () { ...you write this code... }
 public void remove () { ...you write this code... }
}
```



# Code for SparseArrayIterator

```
public class SparseArrayIterator implements Iterator {
 private List current; // pointer to current cell in the list
 SparseArrayIterator(List first) { // the constructor
 current = first;
 }
 public boolean hasNext() {
 return current != null;
 }
 public Object next() {
 Object value = current.value;
 current = current.next
 return value;
 };
 public void remove() {
 // We don't want to implement this, so...
 throw new UnsupportedOperationException();
 }
}
```



# Example, revisited

---

- Instead of:

```
Double max = (Double) array.fetch(0);
for (int i = 0; i < array.length(); i++) {
 Double temp = (Double) array.fetch(i);
 if (temp.compareTo(max) > 0) {
 max = temp;
 }
}
```

- We now need:

```
SparseArrayIterator iterator = new SparseArrayIterator(array);
Double max = (Double) array.fetch(0);
while (iterator.hasNext()) {
 temp = (Double) iterator.next();
 if (temp.compareTo(max) > 0) {
 max = temp;
 }
}
```

- Notice that we use **iterator** in the loop, not **array**



# Not quite there yet...

---

- Our `SparseArrayIterator` is fine for stepping through the elements of an array, but...
  - It doesn't tell us *what index* they were at
  - For some problems, we may need this information
- **Solution #1:** Revise our iterator to tell us, not the value in each list cell, but the index in each list cell
  - Problem: Somewhat more awkward to use, since we would need `array.fetch(iterator.next())` instead of just `iterator.next()`
  - But it's worse than that, because `next` is defined to return an `Object`, so we would have to *wrap* the index
  - We could deal with this by overloading `fetch` to take an `Object` argument
- **Solution #2** (possibly better): Keep `SparseArrayIterator` as is, but *also* write an `IndexIterator`



# IndexIterator

---

- For convenience, we would want **IndexIterator** to return the next index *as an* **int**
- This means that **IndexIterator** *cannot* implement **Iterator**, which defines **next()** to return an **Object**
  - But we can define the same methods (at least those we want)





# Code for IndexIterator

---

```
public class IndexIterator { // does not implement iterator
 private List current; // pointer to current cell in the list
 IndexIterator(List first) { // constructor
 current = first; // just like before
 }
 public boolean hasNext() { // just like before
 return current != null;
 }
 public int next() {
 int index = current.index; // keeps index instead of value
 current = current.next; // just like before
 return index; // returns index instead of value
 }
}
```



# Wrapping the SparseArray class

- If we want a sparse array of, say, **doubles**, we can use the **SparseArray** class by wrapping and unwrapping our values
  - This is a nuisance
  - It's poor style to create another class, say **SparseDoubleArray**, that duplicates all our code
    - Reason: It's much easier and less error-prone if we only have to fix/modify/upgrade our code in one place
  - But we can wrap **SparseArray** itself!



# Code for SparseDoubleArray

- `public class SparseDoubleArray {`
  - `private SparseArray array; // the wrapped array`
  - `public SparseDoubleArray(int size) { // the constructor`  
`array = new SparseArray(size);`  
`}`
  - `// most methods we just "pass on through":`  
`public int length() { return array.length(); }`
  - `// some methods need to do wrapping or unwrapping`  
`public void store(int index, double value) {`  
`array.store(index, new Double(value));`  
`}`
  - `public double fetch(int index) {`  
`Object obj = array.fetch(index);`  
`if (obj == null) return 0.0; // gotta watch out for this case`  
`return ((Double) obj).doubleValue();`  
`}`
  - `// etc.`



# Practical considerations


---

- Writing an ADT such as **SparseArray** can be a lot of work
  - We don't want to duplicate that work for **ints**, for **doubles**, etc.
- If we write **SparseArray** to hold **Objects**, we can use it for anything (including suitably wrapped primitives)
- But—wrappers aren't free
  - A **Double** takes up significantly more space than a **double**
  - Wrapping and unwrapping takes time
- These costs may be acceptable if we don't have a huge number of (non-null) *elements* in our array
  - Note that what is relevant is the number of *actual values*, as opposed to the *defined size* of the array (which is mostly empty)
- Bottom line: Writing a class for **Objects** is usually the simplest and best approach, but sometimes efficiency considerations force you to write a class for a specific type

# Sparse two-dimensional arrays

- Here is an example of a sparse two-dimensional array, and how it can be represented as an *array* of linked lists:

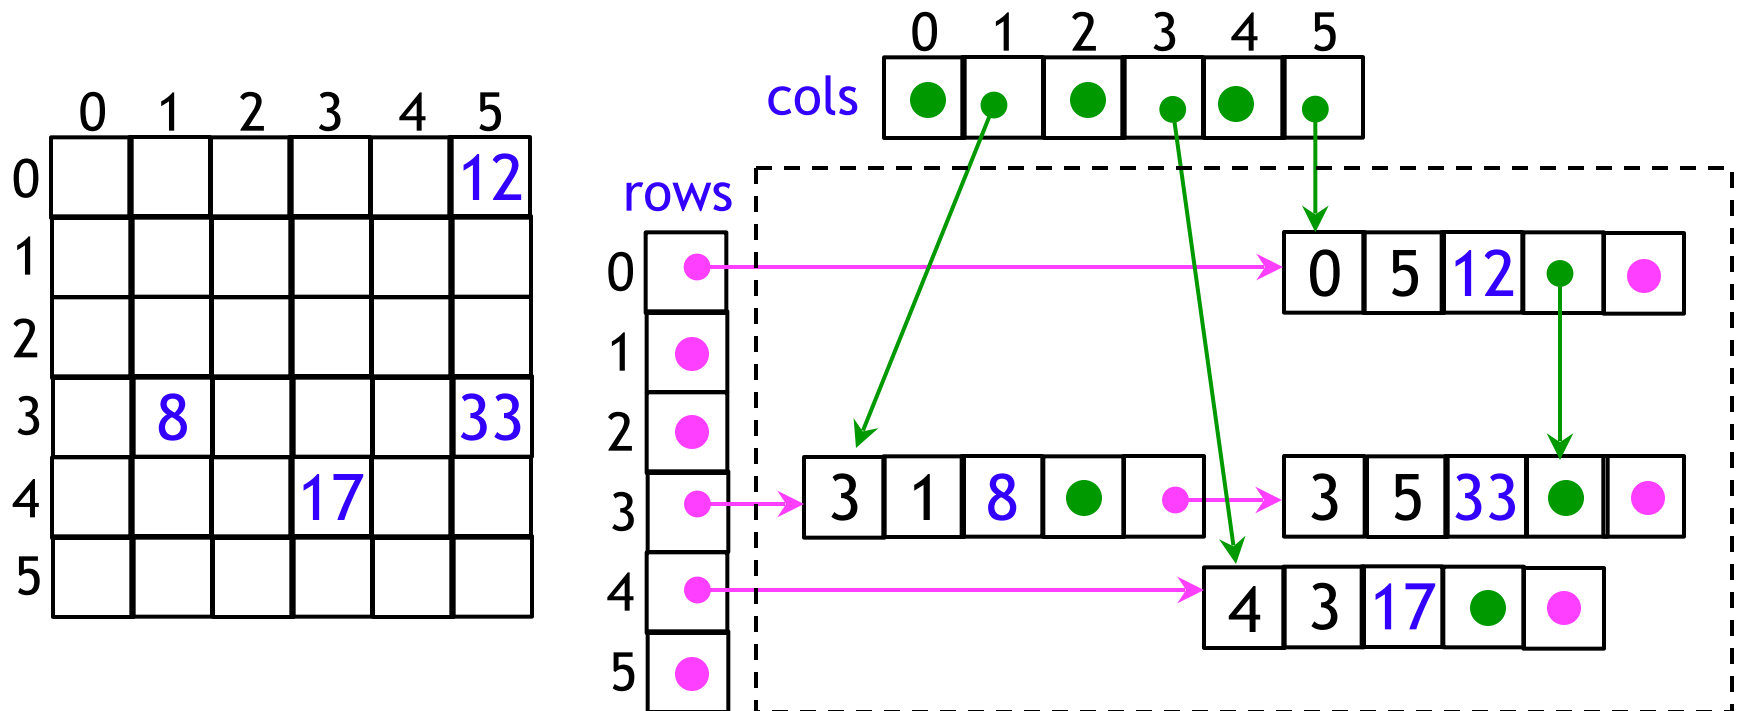
|   | 0 | 1 | 2 | 3  | 4 | 5  |
|---|---|---|---|----|---|----|
| 0 |   |   |   |    |   | 12 |
| 1 |   |   |   |    |   |    |
| 2 |   |   |   |    |   |    |
| 3 |   | 8 |   |    |   | 33 |
| 4 |   |   |   | 17 |   |    |
| 5 |   |   |   |    |   |    |



- With this representation,
  - It is efficient to step through all the elements of a *row*
  - It is expensive to step through all the elements of a *column*
  - Clearly, we could link columns instead of rows
  - Why not both?

# Another implementation

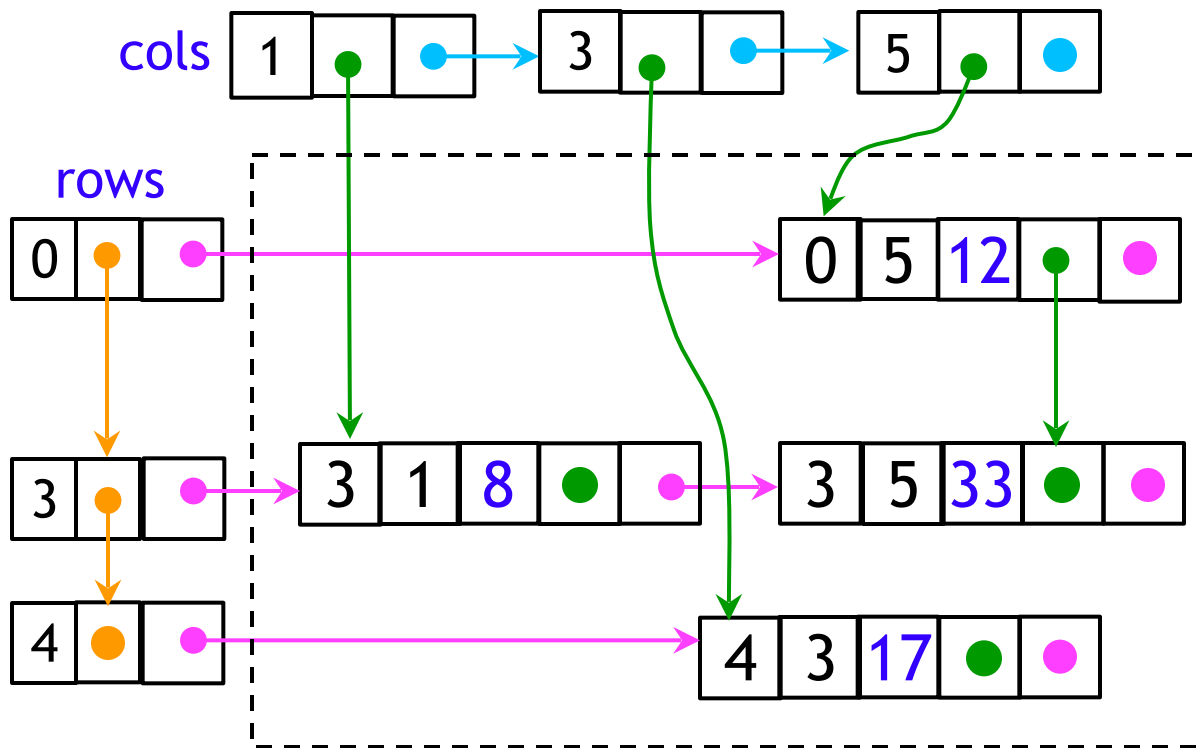
- If we want efficient access to both rows and columns, we need another array and additional data in each node



- Do we really need the row and column number in each node?

# Yet another implementation

- Instead of arrays of pointers to rows and columns, you can use linked lists:



|   | 0 | 1 | 2 | 3  | 4 | 5  |
|---|---|---|---|----|---|----|
| 0 |   |   |   |    |   | 12 |
| 1 |   |   |   |    |   |    |
| 2 |   |   |   |    |   |    |
| 3 |   | 8 |   |    |   | 33 |
| 4 |   |   |   | 17 |   |    |
| 5 |   |   |   |    |   |    |

- Would this be a good data structure for the Penn student grades example?
- This may be the best implementation if most rows and most columns are totally empty



# Considerations

---

- You may need access only by rows, or only by columns
- You may want to access all the elements in a given row without caring what column they are in
  - In this case, you probably should use a Vector instead
- In the most general case, you would want to access by both row and column, just as in an ordinary array





# Looking up an item

---

- The fundamental operation we need is *finding* the item at a given row **i** and column **j**
- Depending on how the array is implemented:
  - We could search a row for a given column number
  - We could search a column for a given row number
    - If we reach a list node with a higher index number, that array location must not be in the linked list
      - If we are doing a **fetch**, we report a value of zero (or **null**)
      - If we are doing a **store**, we may need to insert a cell into the linked list
  - We could choose whether to search by rows or by columns
    - For example you could keep a count of how many elements are in each row and each column (and search the shorter list)



# A sparse 2D array ADT

---

- For a two-dimensional array of Objects, you would need:
  - A constructor:  
`Sparse2DArray(int rows, int columns)`
  - A way to store values in the array:  
`void store(int row, int column, Object value)`
  - A way to get values from the array:  
`Object fetch(int row, int column)`
- Additional useful operations:
  - A way to find the number of rows: `int getRowCount()`
  - A way to find the number of columns: `int getColumnCount()`
  - You may want to find the number of values in each row, or in each column, or in the entire array
  - You almost certainly want row iterators and column iterators



# One final implementation

---

- You could implement a sparse array as a hash table

- For the keys, use something like:

```
class Pair {
 private int row, column;
 Pair (int r, int c) { row = r; column = c; } // constructor
 public boolean equals(Object that) {
 return this.row == that.row && this.column == that.column;
 }
 public int hashCode() {
 return row + column;
 }
}
```


- What are the advantages and disadvantages of this approach?



# Summary

---

- One way to store sparse arrays is as linked lists
- A good ADT provides all the operations a user needs
  - The operations should be *logically complete*
  - They also need to be the *right operations* for *real uses*
- Java **interfaces** provide standardized and (usually) well-thought out skeletons for solving common problems
- It is usually best and most convenient to define ADTs for **Objects** rather than for a specific data type
  - Primitives can be wrapped and used as Objects
  - For even more convenience, the ADT itself can be wrapped
  - The extra convenience also buys us more robust code (because we don't have duplicate almost-the-same copies of our code)
  - Extra convenience comes at a cost in efficiency



# The End

---